

Fused MoE Inference on Blackwell: A Workload-Adaptive Kernel System for DeepSeek-V3

Jiayao Zhang, Jiaoliang Yu
MLSys 2026 FlashInfer-Bench Challenge — Track A

Abstract

We present an agent-assisted, workload-adaptive implementation of the DeepSeek-V3 MoE kernel targeting the NVIDIA B200 (Blackwell, sm_100a) GPU. Our fused Triton pipeline exploits FP8 native tensor core dot products for $2\times$ throughput in GEMM1, a non-atomic two-pass GEMM2-then-reduce architecture, FP16 intermediate buffers with scale-and-cast compensation, and multi-level bucket specialization across BLOCK_M tiers. For the largest traces, a hybrid Triton/CuTe path uses Triton for sparse MoE metadata and epilogues while delegating the grouped GEMM core to CuTe. On 19 real-trace workloads ($T=1$ to $T=14,107$), the system achieves 19/19 correctness, a peak speedup of roughly $56\times$, large- T speedups of roughly $13\text{--}15\times$, and mean speedup around $54\times$ in our reference Modal B200 sessions, as measured by CUPTI GPU-only kernel timing.

1 Introduction

Mixture-of-Experts (MoE) architectures enable massive parameter counts with sub-linear compute growth. DeepSeek-V3 [1] uses 256 routed experts (32 local per device), top-8 routing, hidden dimension 7168, and intermediate dimension 2048—yielding two FP8 block-scaled GEMMs per forward pass.

On NVIDIA B200, naïve per-expert execution wastes hardware: launch overhead dominates small traces, while large traces suffer from atomic scatter in reduction.

This work was developed for Track A of the MLSys 2026 FlashInfer-Bench Challenge [4], where scoring is the sum of CUPTI GPU kernel times across 19 workloads (tolerance atol=1.0, rtol=0.3, ratio=0.9). CUPTI excludes CPU overhead, so optimization is driven by GPU kernel efficiency. Our contributions:

1. A fused Triton [2] pipeline eliminating per-expert launches;
2. Native FP8 tensor core dot in GEMM1, delivering $2\times$ throughput over TF32;

3. Non-atomic GEMM2 with token-centric reduce, yielding $+46\%$ on medium- T ;
4. FP16 intermediate buffers with $\times 0.125/\times 8.0$ compensation (-50% bandwidth);
5. Workload-aware dispatch, including singleton, tight-padding, static-cap, and CuTe-assisted large- T buckets.

2 System Architecture

2.1 Pipeline Overview

Figure 1 shows the generic five-stage pipeline: **(1) Routing:** Triton DeepSeek-V3 no-aux routing (sigmoid, group top-4, global top-8, normalization) with one program per token. **(2) Sorting:** Tile-level parallel histogram with atomic offsets and padded row construction. **(3) GEMM1+SwiGLU:** Fused W_1/W_3 projections in a shared K-loop with `tl.dot(fp8, fp8)`. **(4) GEMM2:** Non-atomic down-projection writing to the `expert_out` buffer. **(5) Token Reduce:** Each output token reads $K=8$ expert contributions, accumulates in FP32, writes bf16.




Figure 1: Generic MoE pipeline. Singleton and large- T workloads reuse the same stages with different kernels.

2.2 Dispatch and Specialization

Workloads span four orders of magnitude, so a single tile shape is suboptimal. We therefore organize execution around a default BLOCK_M tiering rule plus a small number of explicit trace-specific overrides (Table 1). The main overrides are: singleton $T=1$, tight BLOCK_M=8 for fragmented $T=53\text{--}59$, isolated $T=901$, and a hybrid Triton/CuTe path for the largest traces. For large traces,

the dispatcher also switches from static overlunch to exact block-count launch using `total_blocks` produced by sorting.

Table 1: Workload-adaptive dispatch.

Regime	BM	Path
$2 \leq T \leq 128$	16	Generic Triton
$129 \leq T \leq 512$	32	Generic Triton
$513 \leq T < 4096$	64	Generic Triton
$T=1$	16	Singleton Triton
$T=53-59$	8	Tight-padding Triton
$T=901$	64	Isolated Triton
$T=11948, 14107$	128	Triton/CuTe

3 Key Optimizations

3.1 FP8 Native Tensor Core Dot

Blackwell’s `tcgen05.mma` supports native $\text{FP8} \times \text{FP8}$ dot with FP32 accumulation. We apply block-scale dequantization *after* the dot:

```
partial = tl.dot(a, b) # fp8 x fp8
acc += partial * a_scale * b_scale
```

This yields $\sim 2\times$ throughput vs. TF32, which requires $\text{FP8} \rightarrow \text{FP32}$ dequantization before the dot. GEMM1 stores activations and weights in `float8_e4m3fn` with per-128 scaling and fuses W_1/W_3 with inline SwiGLU.

3.2 Non-Atomic GEMM2 + Token-Centric Reduce

Naïve `tl.atomic_add` causes $6\times$ bandwidth amplification (12 B/element FP32 RMW vs. 2 B bf16 store). We therefore split GEMM2 and reduction: GEMM2 writes non-atomic `expert_out` (bf16 where enabled), then a token-centric reduce loads 8 contributions per token. For large T , 2D tiling cuts the reduce grid $8\times$. A/B testing showed this was one of the largest single algorithmic gains, especially on the medium- T workloads where atomic write contention previously dominated.

3.3 FP16 Intermediate with Scale Compensation

The Intermediate buffer [`num_padded`, 2048] dominates HBM traffic. We scale by $\times 0.125$ in GEMM1, store FP16, then compensate by $\times 8.0$ in GEMM2. For $T < 32$, a `constexpr` branch falls back to FP32 to avoid SwiGLU outlier overflow. A/B: mean +4.5% (13/19 improved).

3.4 Workload-Adaptive Bucket Specialization

Column-major dispatch (`GROUP_M=32/64`) improves L2 reuse by mapping consecutive blocks to different tokens of the same expert. The $T=1$ path is a dedicated decode implementation: it bypasses generic histogram/offset metadata and directly fuses routing, packing, GEMM1, GEMM2, and reduction around a fixed small workspace. For fragmented $T=53-59$, `BLOCK_M=8` reduces padding and invalid-row compute. For $T=901$, an isolated path targets the transition regime.

3.5 Routing, Sorting, and Autotuning Support

Several supporting optimizations are smaller individually but important to the final system. Routing uses gather-first and pre-allocated workspaces to reduce overhead; sorting uses tile-parallel histogram/offset construction and, for large traces, can reuse the histogram structure in dispatch. Autotuning remains workload-specific: GEMM1/GEMM2 tile shapes, `BLOCK_K`, and stage depth are selected per regime rather than by a single global configuration.

3.6 CuTe-Assisted Large-T GEMM

For the largest traces, bottlenecks shift from launch/padding overhead to sustained grouped GEMM throughput. We therefore use a hybrid path for $T=11,948$ and $T=14,107$: Triton handles sparse metadata and epilogues, while CuTe executes the grouped GEMM core. We do not use CuTe for small/mid- T , where metadata and setup dominate.

3.7 bf16 Expert Output Buffer

The `expert_out` buffer [`MAX_PADDED`, 7168] is stored in bf16 ($T \geq 32$), saving 50% write bandwidth. bf16’s 3.4×10^{38} range avoids the fp16 overflow that failed on 16/19 workloads. A/B: +0.4% mean, +2.5% at $T=14,107$.

3.8 Compiler Hints and Cache Policy

Two compiler hints yielded a combined +4.9% mean speedup: (1) **Disable LLVM LSR**: Setting `DISABLE_LLVM_OPT="disable-lsr"` avoids extra register pressure and spills in SwiGLU addressing. (2) **Asymmetric L2 eviction**: Expert weights use `evict_last`; token activations use `evict_first`.

Table 2: Failed optimization attempts.

Approach	Result	Root Cause
FP8 Intermediate	0/19	3-bit mantissa; SwiGLU exceeds fp8 range
bf16 Intermediate	8/19	Low mantissa precision; ptxas register failure
Persistent GEMM2	−1.0%	K=2048 too short for weight reuse
TMA Descriptor	−4%	Setup cost not amortized over 16 K-tiles
Warp Specialize	Crash	TTGIR lowering failure; slow fallback
Fused GEMM2+Reduce	−4.5%	fp32 atomic RMW costs 6× bf16 store BW
CuTe for mid- T	No gain	Metadata and host sync not amortized
GEMM2 FP8 Dot	5/19	fp16→fp8 cast remains too lossy

3.9 2D Tiled Token Reduce

The original 1D reduce launched one program per output token, producing >100K programs at $T=14,107$. We switch to 2D [BLOCK_T, BLOCK_N] tiling, cutting the grid >10× and improving ILP. A/B: roughly +2.5–2.9% on the largest traces.

3.10 Deep Software Pipelining for GEMM2

GEMM2’s FP16 A-side incurs high HBM latency. Extending autotuning to `num_stages=6/7/8` overlaps compute with reads and yields +2.1% mean (+10.1% peak) on medium- T workloads.

4 Negative Results

Table 2 summarizes the main rejected experiments.

The failures fall into two groups: lower-precision activation buffers break correctness (Figure 2), while persistence, TMA, fusion, and mid- T CuTe do not amortize setup, metadata, or atomic costs. We retain only changes that pass correctness and exceed the same-session noise floor.

	FP8 (3-bit)	bf16 (7-bit)	FP16 (10-bit)	FP32 (23-bit)
Intermediate (GEMM1 → GEMM2)	0/19 abs>10K	8/19 ptxas fai	19/19 +4.5%	19/19 baseline
expert_out (GEMM2 output)	N/A	19/19 +0.4%	3/19 overflow	19/19 baseline

■ 19/19 PASS
 ■ Partial pass
 ■ FAIL (0/19)

Figure 2: Precision compatibility matrix. Intermediate requires $\geq$ FP16; expert_out tolerates bf16.

5 Evaluation

Setup. We evaluate with the official FlashInfer-Bench CLI on NVIDIA B200 (sm_100a). Scoring uses CUPTI GPU kernel time over 19 real-trace workloads; CPU overhead is excluded. All claims use same-session A/B testing with a $\pm 2\%$ mean noise floor.

5.1 Results

Figure 3 and Table 3 summarize representative performance by workload regime. **Correctness:** 19/19. **Best observed speedup:** $\sim 56\times$. **Mean speedup:** $\sim 54\times$. Small and fragmented traces benefit most because the baseline pays repeated per-expert launch overhead while our fused path amortizes routing, sorting, and GEMM dispatch. $T=1$ is best treated as a singleton specialization rather than part of the grouped-kernel trendline. The two largest traces sustain roughly 13–15× speedup in reference Modal sessions.




Figure 3: Representative speedup by workload regime.

Table 3: Representative benchmark regimes.

T	Traces	Speedup	Bottleneck
1	singleton	specialized	GEMM2
7–80	small/mid traces	$\sim 40\text{--}70\times$	GEMM1
901	isolated	session-dep.	transition
11948–14107	large- T	$\sim 13\text{--}15\times$	GEMM2

5.2 Profiling Analysis

NCU profiling reveals three bottleneck regimes (Figures 4–5): **Small/medium T (7–80):** GEMM1 dominates because expert weights are repeatedly loaded for very few routed rows, with substantial padding waste. **$T=901$:** the workload sits in a GEMM1/GEMM2 transition regime, motivating an isolated path. **Large hybrid T (11,948–14,107):** GEMM2 becomes the main bottleneck, while CuTe improves the regular grouped-GEMM core. Token reduce remains lightweight across regimes.

5.3 Optimization Impact

Figure 6 shows cumulative impact over 15 rounds. The largest jumps come from eliminating per-expert




Figure 4: Representative kernel-time breakdown by workload regime.




Figure 5: Representative effective TFLOPS vs. token count.

launches, building the FP8-dot/non-atomic GEMM2 backbone, and then adding workload-aware specialization. Later gains are smaller but cumulative: compiler hints, deeper GEMM2 pipelining, bf16 expert_out, exact large- T dispatch, isolated runtimes, and tight mid- T BLOCK_M tuning.

5.4 Measurement Noise

Modal B200 shared instances exhibit $\pm 2\%$ mean noise and $\pm 15\%$ per-workload noise for small T (where kernel times are < 0.2 ms and hardware jitter dominates), but $< 1\%$ for large T . Cross-session drift reaches $\pm 20\text{--}30\%$ due to co-tenant interference; all claims use same-session A/B testing. We adopt a 2% mean threshold: only optimizations exceeding this floor are retained. The contest tolerance (atol=1.0, rtol=0.3, 90% element ratio) is $\sim 100\times$ wider than our internal checks (atol=0.01), enabling bf16 expert output that would otherwise fail strict validation.

6 Human vs. Agent Contributions

Human contributors led the profiling-driven strategy: they identified workload-adaptive dispatch as the organizing principle, decided when pure Triton had reached diminishing returns, proposed introducing CuTe DSL for the largest grouped-GEMM traces, and selected which experiments to keep or revert. They also interpreted benchmark noise, made final correctness/risk decisions, and maintained the contest-facing solution. The agent assisted by inspecting profiler traces, expanding human hypotheses into bounded Triton/CuTe experiments, com-




Figure 6: Representative cumulative speedup across major optimization phases.

paring A/B results, and summarizing failed directions. All final optimizations were human-reviewed against correctness, benchmark stability, and simplicity before inclusion.

7 Conclusion

We presented a workload-adaptive MoE kernel system achieving up to roughly $56\times$ speedup and roughly $54\times$ mean speedup on B200 reference Modal sessions through profiling-driven optimization over 15 rounds. The final system combines FP8 dot, non-atomic GEMM2, FP16 intermediate, workload-aware dispatch, CuTe-assisted large- T GEMM, 2D reduce, and deeper GEMM2 pipelining. Over the same period, we rejected persistent kernels, TMA, warp specialization, fused reduce, and CuTe for mid- T —all cases where setup cost, synchronization, or $K=2048$ (only 16 K-loop iterations) prevented amortization.

Future work. The remaining gap lies in small- T memory-boundedness (5.6% of FP8 peak at $T=7$). L2 persistent cache and epsilon expert skipping remain disabled pending further tuning. `tl.dot_scaled` with MXFP8 [3] achieved only `matched_ratio=0.17–0.32`, confirming shared exponents are too coarse for GEMM2. As MoE models scale, workload-adaptive kernels will be essential.

References

- [1] DeepSeek-AI. DeepSeek-V3 technical report. *arXiv:2412.19437*, 2024.
- [2] P. Tillet, H. T. Kung, D. Cox. Triton: An intermediate language and compiler for tiled neural network computations. *MLSys*, 2019.
- [3] B. Darvish Rouhani et al. Microscaling data formats for deep learning. *arXiv:2310.10537*, 2023.
- [4] Z. Ye et al. FlashInfer: Efficient and customizable attention engine for LLM inference serving. *arXiv:2501.01005*, 2025.