

Fused MoE Inference on Blackwell: A Workload-Adaptive Kernel System for DeepSeek-V3

Jiayao Zhang, Jiaoliang Yu
MLSys 2026 FlashInfer-Bench Challenge — Track A

Abstract

We present an agent-assisted, workload-adaptive implementation of the DeepSeek-V3 MoE kernel targeting the NVIDIA B200 (Blackwell, sm_100a) GPU. Our fused Triton pipeline exploits FP8 native tensor core dot products for $2\times$ throughput in GEMM1, a non-atomic two-pass GEMM2-then-reduce architecture, FP16 intermediate buffers with scale-and-cast compensation, and multi-level bucket specialization across BLOCK_M tiers. For the largest traces, a hybrid Triton/CuTe path uses Triton for sparse MoE metadata and epilogues while delegating the grouped GEMM core to CuTe. On 19 real-trace workloads ($T=1$ to $T=14,107$), the system achieves 19/19 correctness with a peak speedup of $\sim 56\times$, large- T speedups of $\sim 13\text{--}15\times$, and a mean of $\sim 54\times$ on reference Modal B200 sessions, as measured by CUPTI GPU-only kernel timing.

1 Introduction

Mixture-of-Experts (MoE) architectures enable massive parameter counts with sub-linear compute growth. DeepSeek-V3 [1] uses 256 routed experts (32 local per device), top-8 routing, hidden dimension 7168, and intermediate dimension 2048—yielding two FP8 block-scaled GEMMs per forward pass.

On NVIDIA B200, naïve per-expert execution wastes hardware: launch overhead dominates small traces, while large traces suffer from atomic scatter in reduction.

This work was developed for Track A of the MLSys 2026 FlashInfer-Bench Challenge [4], where scoring is the sum of CUPTI GPU kernel times across 19 workloads (tolerance atol=1.0, rtol=0.3, ratio=0.9). CUPTI excludes CPU overhead, so optimization is driven by GPU kernel efficiency. Our contributions:

1. A fused Triton [2] pipeline eliminating per-expert launches;

2. Native FP8 tensor core dot in GEMM1, delivering $2\times$ throughput over TF32;
3. Non-atomic GEMM2 with token-centric reduce, yielding +46% on medium- T ;
4. FP16 intermediate buffers with $\times 0.125/\times 8.0$ compensation (-50% bandwidth);
5. Workload-aware dispatch, including singleton, tight-padding, static-cap, and CuTe-assisted large- T buckets.

2 System Architecture

2.1 Pipeline Overview

Figure 1 shows the generic five-stage pipeline: (1) Routing: Triton DeepSeek-V3 no-aux routing (sigmoid, group top-4, global top-8, normalization) with one program per token. (2) Sorting: Tile-level parallel histogram with atomic offsets and padded row construction. (3) GEMM1+SwiGLU: Fused W_1/W_3 projections in a shared K-loop with `tl.dot(fp8, fp8)`. (4) GEMM2: Non-atomic down-projection writing to the `expert_out` buffer. (5) Token Reduce: Each output token reads $K=8$ expert contributions, accumulates in FP32, writes bf16.




Figure 1: Generic MoE pipeline. Singleton and large- T workloads reuse the same stages with different kernels.

2.2 Dispatch and Specialization

Workloads span four orders of magnitude, so a single tile shape is suboptimal. We therefore organize execution around a default BLOCK_M tiering rule plus a small number of explicit trace-specific overrides (Table 1). The main overrides are: singleton $T=1$,

tight BLOCK_M=8 for fragmented $T=53\text{--}59$, isolated $T=901$, and a hybrid Triton/CuTe path for the largest traces. For $T \geq 4096$, the dispatcher switches from static overlaunch to exact block-count launch using total_blocks produced by sorting, avoiding $\sim 30\%$ idle blocks.

Table 1: Workload-adaptive dispatch.

Regime	BM	Path
$2 \leq T \leq 128$	16	Generic Triton
$129 \leq T \leq 512$	32	Generic Triton
$513 \leq T \leq 2047$	64	Generic Triton
$T \geq 2048$	128	Generic Triton (exact launch)
$T=1$	16	Singleton Triton
$T=53\text{--}59$	8	Tight-padding Triton
$T=901$	64	Isolated Triton
$T=11948, 14107$	128	Triton/CuTe hybrid

3 Key Optimizations

3.1 FP8 Native Tensor Core Dot

Blackwell’s tcgen05.mma supports native FP8 $\times$ FP8 dot with FP32 accumulation. We apply block-scale dequantization after the dot:

```
partial = tl.dot(a, b) # fp8 x fp8
acc += partial * a_scale * b_scale
```

This yields $\sim 2\times$ throughput vs. TF32, which requires FP8 $\rightarrow$ FP32 dequantization before the dot. GEMM1 stores activations and weights in float8_e4m3fn with per-128 scaling and fuses W_1/W_3 with inline SwiGLU.

3.2 Non-Atomic GEMM2 + Token-Centric Reduce

Naïve tl.atomic_add causes $6\times$ bandwidth amplification (12 B/element FP32 RMW vs. 2 B bf16 store). We split GEMM2 and reduction: GEMM2 writes non-atomic expert_out (bf16 where enabled), then a token-centric reduce loads 8 contributions per token. For large T , 2D tiling cuts the reduce grid $>10\times$. A/B testing confirmed this as one of the largest single gains, especially on medium- T workloads where atomic contention previously dominated.

3.3 FP16 Intermediate with Scale Compensation

The Intermediate buffer [num_padded, 2048] dominates HBM traffic. We scale by $\times 0.125$ in GEMM1, store FP16, then compensate by $\times 8.0$ in GEMM2.

For $T < 32$, a constexpr branch falls back to FP32 to avoid SwiGLU outlier overflow. A/B: mean +4.5% (13/19 improved).

3.4 Workload-Adaptive Bucket Specialization

Column-major dispatch (GROUP_M=32/64) improves L2 reuse by mapping consecutive blocks to different tokens of the same expert. The $T=1$ path is a dedicated decode implementation: it bypasses generic histogram/offset metadata and uses two fused kernels (GEMM1+SwiGLU and GEMM2+reduce) around a fixed small workspace. For fragmented $T=53\text{--}59$, BLOCK_M=8 reduces padding and invalid-row compute. For $T=901$, an isolated path targets the transition regime.

3.5 Routing, Sorting, and Autotuning Support

Several supporting optimizations are smaller individually but important to the final system. Routing uses a Triton kernel implementing group-filter then global top-K selection with pre-allocated workspaces; sorting uses tile-parallel histogram/offset construction and, for large traces, reuses the histogram structure in dispatch. Autotuning remains workload-specific: GEMM1/GEMM2 tile shapes, BLOCK_K, and stage depth are selected per regime rather than by a single global configuration.

3.6 CuTe-Assisted Large-T GEMM

For the largest traces, bottlenecks shift from launch/padding overhead to sustained grouped GEMM throughput. We therefore use a hybrid path for $T=11,948$ and $T=14,107$: Triton handles sparse metadata and epilogues, while CuTe [5] executes the grouped GEMM core. We do not use CuTe for small/mid- T , where metadata and setup dominate.

3.7 bf16 Expert Output Buffer

The expert_out buffer [MAX_PADDED, 7168] is stored in bf16 ($T \geq 32$), saving 50% write bandwidth. bf16’s 3.4×10^{38} range avoids the fp16 overflow that caused inf on 3/19 workloads. A/B: +0.4% mean, +2.5% at $T=14,107$.

3.8 Compiler Hints and Cache Policy

Two compiler hints yielded a combined +4.9% mean speedup: (1) Disable LLVM LSR: Setting

Table 2: Failed optimization attempts.

Approach	Result	Root Cause
FP8 Intermediate	0/19	3-bit mantissa; SwiGLU exceeds fp8 range
bf16 Intermediate	8/19	Low mantissa precision; ptxas register failure
Persistent GEMM2	−1.0%	K=2048 too short for weight reuse
TMA Descriptor	−4%	Setup cost not amortized over 16 K-tiles
Warp Specialize	Crash	TTGIR lowering failure; slow fallback
Fused GEMM2+Reduce	−4.5%	fp32 atomic RMW costs 6× bf16 store BW
CuTe for mid- T	No gain	Metadata and host sync not amortized
GEMM2 FP8 Dot	5/19	fp16→fp8 cast remains too lossy

DISABLE_LLVM_OPT="disable-lsr" avoids extra register pressure and spills in SwiGLU addressing. (2) Asymmetric L2 eviction: Expert weights use evict_last; token activations use evict_first.

3.9 2D Tiled Token Reduce

The original 1D reduce launched one program per (token, N-block) pair, producing a large grid at $T=14,107$. Switching to 2D [BLOCK_T, BLOCK_N] tiling amortizes the per-token loop overhead across BLOCK_T rows, cutting the grid $>10\times$ and improving occupancy. A/B: +2.5–2.9% on the largest traces.

3.10 Deep Software Pipelining for GEMM2

GEMM2’s FP16 A-side incurs high HBM latency. Extending autotuning to num_stages=6/7/8 overlaps compute with reads and yields +2.1% mean (+10.1% peak) on medium- T workloads.

4 Negative Results

Table 2 summarizes the main rejected experiments. The failures fall into two groups: lower-precision activation buffers break correctness (Figure 2), while persistence, TMA, fusion, and mid- T CuTe do not amortize setup, metadata, or atomic costs. We retain only changes that pass correctness and exceed the same-session noise floor.




Figure 2: Precision compatibility matrix. Intermediate requires $\geq$ FP16; expert_out tolerates bf16.

5 Evaluation

Setup. We evaluate with the official FlashInfer-Bench CLI on NVIDIA B200 (sm_100a). Scoring uses CUPTI GPU kernel time over 19 real-trace workloads; CPU overhead is excluded. All claims use same-session A/B testing with a $\pm 2\%$ mean noise floor.

5.1 Results

Figure 3 and Table 3 summarize representative performance by workload regime. Correctness: 19/19. Best observed speedup: $\sim 56\times$. Mean speedup: $\sim 54\times$. Small and fragmented traces benefit most because the baseline pays repeated per-expert launch overhead while our fused path amortizes routing, sorting, and GEMM dispatch. $T=1$ is best treated as a singleton specialization rather than part of the grouped-kernel trendline. The two largest traces sustain $\sim 13\text{--}15\times$ speedup in reference Modal sessions.




Figure 3: Representative speedup by workload regime.

Table 3: Representative benchmark regimes.

T	Traces	Speedup	Bottleneck
1	singleton decode	dedicated	GEMM2
7–80	small/mid traces	$\sim 46\text{--}56\times$	GEMM1
901	transition regime	$\sim 18\times$	GEMM2
11948–14107	large- T	$\sim 13\text{--}15\times$	GEMM2

5.2 Profiling Analysis

NCU profiling reveals three bottleneck regimes (Figures 4–5): Small/medium T (7–80): GEMM1 dominates because expert weights are repeatedly




Figure 4: Representative kernel-time breakdown by workload regime.




Figure 5: Representative effective TFLOPS vs. token count.

loaded for very few routed rows, with substantial padding waste. $T=901$: the workload sits in a GEMM1/GEMM2 transition regime, motivating an isolated path. Large hybrid T (11,948–14,107): GEMM1 still takes the larger time share (2× the FLOPs from fused W_1/W_3), but GEMM2 is the efficiency bottleneck—its FP16 A-side prevents FP8 tensor core use, capping throughput well below peak. CuTe improves the grouped-GEMM core. Token reduce remains lightweight across regimes.

5.3 Optimization Impact

Figure 6 shows cumulative impact over 23 rounds of optimization. The largest jumps come from eliminating per-expert launches, building the FP8-dot/non-atomic GEMM2 backbone, and then adding workload-aware specialization. Later gains are smaller but cumulative: compiler hints, deeper GEMM2 pipelining, bf16 expert_out, exact large- T dispatch, isolated runtimes, and tight mid- T BLOCK_M tuning.

5.4 Measurement Noise

Modal B200 shared instances exhibit $\pm 2\%$ mean noise and $\pm 15\%$ per-workload noise for small T , but $< 1\%$ for large T . Cross-session drift reaches $\pm 20-30\%$; all claims use same-session A/B testing with a 2% mean threshold. The contest tolerance (atol=1.0, rtol=0.3, 90%) is $\sim 100\times$ wider than our internal checks, enabling bf16 expert output.




Figure 6: Representative cumulative speedup across major optimization phases.

6 Tools & Languages

All kernels use Triton [2], chosen for rapid iteration, built-in autotuning, and first-class FP8 tensor core support on Blackwell. For the two largest traces, the grouped-GEMM core is offloaded to CuTe DSL [5] for tighter warp-level control. Host orchestration uses PyTorch. We evaluated but did not adopt raw CUDA, Tilelang, or CuTile: Triton’s autotuner and FP8 intrinsics sufficed for 17/19 workloads, and CuTe covered the remaining two.

7 Human vs. Agent Contributions

Human contributors led the profiling-driven strategy: they identified workload-adaptive dispatch as the organizing principle, decided when pure Triton had reached diminishing returns, proposed introducing CuTe DSL for the largest grouped-GEMM traces, and selected which experiments to keep or revert. They also interpreted benchmark noise, made final correctness/risk decisions, and maintained the contest-facing solution. The agent assisted by inspecting profiler traces, expanding human hypotheses into bounded Triton/CuTe experiments, comparing A/B results, and summarizing failed directions. All final optimizations were human-reviewed against correctness, benchmark stability, and simplicity before inclusion.

8 Conclusion

We presented a workload-adaptive MoE kernel system achieving $\sim 56\times$ peak and $\sim 54\times$ mean speedup on B200 through 23 rounds of profiling-driven optimization, combining FP8 dot, non-atomic GEMM2, FP16 intermediate, workload-aware dispatch, and CuTe-assisted large- T GEMM. Eight rejected approaches fall into precision failures and amortization failures where $K=2048$ does not justify setup cost.

References

- [1] DeepSeek-AI. DeepSeek-V3 technical report. arXiv:2412.19437, 2024.
- [2] P. Tillet, H. T. Kung, D. Cox. Triton: An intermediate language and compiler for tiled neural network computations. MLSys, 2019.
- [3] B. Darvish Rouhani et al. Microscaling data formats for deep learning. arXiv:2310.10537, 2023.
- [4] Z. Ye et al. FlashInfer: Efficient and customizable attention engine for LLM inference serving. arXiv:2501.01005, 2025.
- [5] A. Kerr et al. CUTLASS: Fast linear algebra in CUDA C++. NVIDIA, 2017–2025.