

Fused MoE Inference on Blackwell: A Workload-Adaptive Kernel System for DeepSeek-V3

Jiayao Zhang, Jiaoliang Yu
MLSys 2026 FlashInfer-Bench Challenge — Track A

Abstract

We present a workload-adaptive implementation of the DeepSeek-V3 MoE kernel targeting the NVIDIA B200 (Blackwell, sm_100a) GPU. Our fused Triton pipeline exploits FP8 native tensor core dot products for $2\times$ throughput in GEMM1, a non-atomic two-pass GEMM2-then-reduce architecture, FP16 intermediate buffers with scale-and-cast compensation, and multi-level bucket specialization across BLOCK_M tiers. For the largest traces, a hybrid Triton/CuTe path uses Triton for sparse MoE metadata and epilogues while delegating the dense grouped GEMM core to CuTe. On 19 real-trace workloads ($T=1$ to $T=14,107$), the system achieves 19/19 correctness, a peak speedup of roughly $56\times$, large- T speedups of roughly $13\text{--}15\times$, and mean speedup around $54\times$ in our reference Modal B200 sessions, as measured by CUPTI GPU-only kernel timing.

1 Introduction

Mixture-of-Experts (MoE) architectures enable massive parameter counts with sub-linear compute growth [1]. DeepSeek-V3 [2] uses 256 routed experts (32 local per device), top-8 routing, hidden dimension 7168, and intermediate dimension 2048—yielding two FP8 block-scaled GEMMs per forward pass.

On NVIDIA B200, naïve per-expert execution wastes hardware: launch overhead dominates small traces, while large traces suffer from atomic scatter in reduction.

This work was developed for Track A of the MLSys 2026 FlashInfer-Bench Challenge [7], where scoring is the sum of CUPTI GPU kernel times across 19 workloads (tolerance $\text{atol}=1.0$, $\text{rtol}=0.3$, $\text{ratio}=0.9$). CUPTI excludes CPU overhead, so optimization is driven by GPU kernel efficiency. Our contributions:

1. A fused Triton [3] pipeline eliminating per-expert launches;
2. Native FP8 tensor core dot in GEMM1, delivering $2\times$ throughput over TF32;
3. Non-atomic GEMM2 with token-centric reduce, yielding +46% on medium- T ;
4. FP16 intermediate buffers with $\times 0.125/\times 8.0$ compensation (-50% bandwidth);
5. Workload-aware dispatch, including singleton, tight-padding, static-cap, and CuTe-assisted large- T buckets.




Figure 1: Generic MoE pipeline. Singleton and large- T workloads reuse the same stages with different kernels.

2 System Architecture

2.1 Pipeline Overview

Figure 1 shows the generic five-stage pipeline: **(1) Routing:** Triton DeepSeek-V3 no-aux routing (sigmoid, group top-4, global top-8, normalization) with one program per token. **(2) Sorting:** Tile-level parallel histogram with atomic offsets; very small traces may use a hybrid CPU path to replace many tiny GPU launches in wall time. **(3) GEMM1+SwiGLU:** Fused W_1/W_3 projections in a shared K-loop with `tl.dot(fp8, fp8)`. **(4) GEMM2:** Non-atomic down-projection writing to `expert_out` buffer. **(5) Token Reduce:** Each output token reads $K=8$ expert contributions, accumulates in FP32, writes bf16.

2.2 Dispatch and Specialization

Workloads span four orders of magnitude, so a single tile shape is suboptimal. We therefore organize execution around a default BLOCK_M tiering rule plus a small number of explicit trace-specific overrides (Table 1). The main overrides are: singleton $T=1$, tight BLOCK_M=8 for fragmented $T=53\text{--}59$, isolated $T=901$, and a hybrid Triton/CuTe path for the largest traces. For large traces, the dispatcher also switches from static overlaunch to exact block-count launch using `total_blocks` produced by sorting.

3 Key Optimizations

3.1 FP8 Native Tensor Core Dot

Blackwell’s `tcgen05.mma` supports native $\text{FP8}\times\text{FP8}$ dot with FP32 accumulation. We apply block-scale dequantization *after* the dot:

```

partial = tl.dot(a, b) # fp8 x fp8
acc += partial * a_scale * b_scale

```

Table 1: Workload-adaptive dispatch: generic Triton tiers plus explicit overrides.

Regime	BLOCK_M	Specialization
Default: $2 \leq T \leq 128$	16	Generic Triton small/mid- T tier
Default: $129 \leq T \leq 512$	32	Generic Triton medium tier
Default: $513 \leq T < 4096$	64	Generic Triton large tier
Override: $T=1$	16	Singleton fused Triton path
Override: $T \in \{53, \dots, 59\}$	8	Tight-padding mid- T path
Override: $T=901$	64	Isolated static-cap Triton path
Override: $T \in \{11948, 14107\}$	128	Hybrid Triton/CuTe path

This yields $\sim 2\times$ throughput vs. TF32, which requires FP8 $\rightarrow$ FP32 dequantization before the dot. GEMM1 stores activations and weights in `float8_e4m3fn` with per-128 scaling and fuses W_1/W_3 with inline SwiGLU.

3.2 Non-Atomic GEMM2 + Token-Centric Reduce

Naïve `tl.atomic_add` causes $6\times$ bandwidth amplification (12 B/ele FP32 RMW vs. 2 B bf16 store). We therefore split GEMM2 and reduction: GEMM2 writes bf16 `expert_out`, then a token-centric reduce loads 8 contributions per token. For large T , 2D tiling cuts the reduce grid $8\times$. A/B testing showed this was one of the largest single algorithmic gains, especially on the medium- T workloads where atomic write contention previously dominated.

3.3 FP16 Intermediate with Scale Compensation

The Intermediate buffer [`num_padded`, 2048] dominates HBM traffic. We scale by $\times 0.125$ in GEMM1, store FP16, then compensate by $\times 8.0$ in GEMM2. For $T < 32$, a `constexpr` branch falls back to FP32 to avoid SwiGLU outlier overflow. A/B: mean +4.5% (13/19 improved).

3.4 Workload-Adaptive Bucket Specialization

Column-major dispatch (`GROUP_M=32/64`) improves L2 reuse by mapping consecutive blocks to different tokens of the same expert. The $T=1$ path is a dedicated decode implementation: it bypasses generic histogram/offset metadata and directly fuses routing, packing, GEMM1, GEMM2, and reduction around a fixed small workspace. For fragmented $T=53\text{--}59$, `BLOCK_M=8` reduces padding and invalid-row compute. For $T=901$, an isolated path targets the transition regime.

3.5 Routing, Sorting, and Autotuning Support

Several supporting optimizations are smaller individually but important to the final system. Routing uses gather-first and pre-allocated workspaces to reduce overhead; sorting uses tile-parallel histogram/offset construction and, for large traces, can reuse the histogram structure in dispatch. Autotuning remains workload-specific: GEMM1/GEMM2 tile shapes, `BLOCK_K`, and stage depth are selected per regime rather than by a single global configuration.

3.6 CuTe-Assisted Large-T GEMM

For the largest traces, bottlenecks shift from launch/padding overhead to sustained grouped GEMM throughput. We therefore use a hybrid path for $T=11,948$ and $T=14,107$: Triton handles sparse metadata and epilogues, while CuTe executes the dense grouped GEMM core. We do not use CuTe for small/mid- T , where metadata and setup dominate.

3.7 bf16 Expert Output Buffer

The `expert_out` buffer [`MAX_PADDED`, 7168] is stored in bf16 ($T \geq 32$), saving 50% write bandwidth. bf16’s 3.4×10^{38} range avoids the fp16 overflow that failed on 16/19 workloads. A/B: +0.4% mean, +2.5% at $T=14,107$.

3.8 Compiler Hints and Cache Policy

Two compiler hints yielded a combined +4.9% mean speedup: (1) **Disable LLVM LSR**: Setting `DISABLE_LLVM_OPT="disable-lsr"` avoids extra register pressure and spills in SwiGLU addressing. (2) **Asymmetric L2 eviction**: Expert weights use `evict_last`; token activations use `evict_first`.

3.9 2D Tiled Token Reduce

The original 1D reduce launched one program per output token, producing $>100K$ programs at $T=14,107$. We switch to 2D [`BLOCK_T`, `BLOCK_N`] tiling, cutting the grid $>10\times$ and improving ILP. A/B: roughly +2.5–2.9% on the largest traces.

3.10 Deep Software Pipelining for GEMM2

GEMM2’s FP16 A-side incurs high HBM latency. Extending autotuning to `num_stages=6/7/8` overlaps compute with reads and yields +2.1% mean (+10.1% peak) on medium- T workloads.

4 Negative Results

Table 2 summarizes major failed optimizations explored over 15 rounds.

Precision hierarchy (Figure 2). Intermediate requires $\geq$ FP16; bf16 was tested three times and failed 8–13 workloads. `expert_out` tolerates bf16, while fp16 overflows. FP8 is insufficient for any activation buffer. **Architectural**

Table 2: Failed optimization attempts.

Approach	Result	Root Cause
FP8 Intermediate	0/19	3-bit mantissa; SwiGLU exceeds fp8 max (448)
bf16 Intermediate	8/19	7-bit mantissa insufficient; ptxas register failure
Persistent GEMM2	−1.0%	K=2048 too short for weight tile reuse
TMA Descriptor	−4%	Setup cost not amortized over 16 K-loop iters
Warp Specialize	Crash	TTGIR lowering failure; −34% on simple GEMM
Fused GEMM2+Reduce	−4.5%	fp32 atomic RMW = 6× BW vs bf16 store
CuTe for mid- T	Neutral/−	metadata and host sync not amortized at small expert counts
GEMM2 FP8 Dot	5/19	fp16→fp8 cast: abs=500K direct, 14K scaled

	FP8 (3-bit)	bf16 (7-bit)	FP16 (10-bit)	FP32 (23-bit)
Intermediate (GEMM1 → GEMM2)	0/19 abs>10K txas fail	8/19 txas fail	19/19 +4.5%	19/19 baseline
expert_out (GEMM2 output)	N/A	19/19 +0.4%	3/19 overflow	19/19 baseline

■ 19/19 PASS
 ■ Partial pass
 ■ FAIL (0/19)

Figure 2: Precision compatibility matrix. Intermediate requires $\geq$ FP16; expert_out tolerates bf16.

dead-ends. Persistent kernels, TMA, and warp specialization fail because $K=2048$ is too short to amortize setup. CuTe follows the same rule and only helps at large T .

5 Evaluation

Setup. NVIDIA B200 (sm_100a), PyTorch 2.12.0+cu132, Triton 3.6.0. Scoring: CUPTI GPU kernel time (CPU overhead excluded). 19 real-trace workloads, 15 warmup + 80 measurement iterations. All claims validated by same-session A/B testing ($\pm 2\%$ mean noise floor).

5.1 Results

Figure 3 and Table 3 summarize representative performance by workload regime. **Correctness:** 19/19. **Best observed speedup:** $\sim 56\times$. **Mean speedup:** $\sim 54\times$. Small and fragmented traces benefit most because the baseline pays repeated per-expert launch overhead while our fused path amortizes routing, sorting, and GEMM dispatch. $T=1$ is best treated as a singleton specialization rather than part of the grouped-kernel trendline. The two largest traces sustain roughly 13–




Figure 3: Representative speedup by workload regime.

Table 3: Representative benchmark summary by workload regime.

T regime	Representative traces	Speedup	Bottleneck
$T=1$	singleton decode	specialized path	GEMM2
$T=7-80$	7, 14, 15, 16, 32, 52–62, 80	$\sim 40-70\times$	GEMM1
$T=901$	isolated Triton path	session-dependent	GEMM1
$T=11948-14107$	CuTe-assisted large- T	$\sim 13-15\times$	GEMM2

15 $\times$ speedup in reference Modal sessions.

5.2 Profiling Analysis

NCU profiling reveals three bottleneck regimes (Figures 4–5): **Small/medium T (7–80):** GEMM1 dominates (51–61%), achieving only 103–182 TFLOPS—severely memory-bound due to expert weight cold misses ($T=7$ distributes <1 token per expert yet must load full weight tiles). Padding waste is substantial: at $T=7$, 7 local rows pad to 112 ($16\times$ waste); at $T=52$, 26 rows pad to 272 ($10.5\times$). **Isolated $T=901$:** this workload sits in a transition regime where GEMM1 remains the largest component (52.6%), GEMM2 contributes 35.3%, and routing+sorting adds 7.9%. **Large hybrid T (11,948–14,107):** GEMM2 dominates (49–50%), while GEMM1 drops to roughly 40%; GEMM1 reaches 485–535 TFLOPS and GEMM2 about 216 TFLOPS in the reference sessions. Token reduce remains lightweight: effectively fused for $T=1$, a few percent for small- T , and about 3% on the largest traces.

5.3 Optimization Impact

Figure 6 shows cumulative impact over 15 rounds. The dominant jumps come from three phases: eliminating per-expert launches with monolithic Triton, building the core FP8-dot + non-atomic GEMM2 stack, and then layering workload-aware specialization on top. The non-atomic GEMM2 discovery overturned a prior conclusion that GEMM2 was compute-saturated—the real bottleneck was atomic write contention, where FP32 read-modify-write inflated bandwidth by $6\times$ vs. bf16 non-atomic store. Later gains are smaller but still meaningful: compiler hints, deeper GEMM2 pipelining, bf16 expert_out, exact large- T dispatch, isolated runtimes, and tight mid- T BLOCK_M tuning.




Figure 4: Representative kernel-time breakdown by workload regime.




Figure 5: Representative effective TFLOPS vs. token count.

5.4 Measurement Noise

Modal B200 shared instances exhibit $\pm 2\%$ mean noise and $\pm 15\%$ per-workload noise for small T (where kernel times are < 0.2 ms and hardware jitter dominates), but $< 1\%$ for large T . Cross-session drift reaches ± 20 – 30% due to co-tenant interference; all claims use same-session A/B testing. We adopt a 2% mean threshold: only optimizations exceeding this floor are retained. The contest tolerance ($\text{atol}=1.0$, $\text{rtol}=0.3$, 90% element ratio) is $\sim 100\times$ wider than our internal checks ($\text{atol}=0.01$), enabling bf16 expert output that would otherwise fail strict validation.

6 Conclusion

We presented a workload-adaptive MoE kernel system achieving up to roughly $56\times$ speedup and roughly $54\times$ mean speedup on B200 reference Modal sessions through profiling-driven optimization over 15 rounds. The final system combines FP8 dot, non-atomic GEMM2, FP16 intermediate, workload-aware dispatch, CuTe-assisted large- T GEMM, 2D reduce, and deeper GEMM2 pipelining. Over the same period, we rejected persistent kernels, TMA, warp specialization, fused reduce, and CuTe for mid- T —all cases where setup cost, synchronization, or $K=2048$ (only 16 K-loop iterations) prevented amortization.

Future work. The remaining gap lies in small- T memory-boundedness (5.6% of FP8 peak at $T=7$). L2 persistent cache and epsilon expert skipping remain disabled pending further tuning. `tl.dot_scaled` with MXFP8 [5] achieved only `matched_ratio=0.17–0.32`, confirming shared exponents are too coarse for GEMM2. As MoE models scale [2], workload-adaptive kernels—a direction SGLang [8], MegaBlocks [4], and vLLM [6] explore—will be essential.




Figure 6: Representative cumulative speedup across major optimization phases.

References

- [1] N. Shazeer et al. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. *ICLR*, 2017.
- [2] DeepSeek-AI. DeepSeek-V3 technical report. *arXiv:2412.19437*, 2024.
- [3] P. Tillet, H. T. Kung, D. Cox. Triton: An intermediate language and compiler for tiled neural network computations. *MLSys*, 2019.
- [4] T. Gale, D. Narayanan, C. Young, M. Zaharia. MegaBlocks: Efficient sparse training with mixture-of-experts. *MLSys*, 2023.
- [5] B. Darvish Rouhani et al. Microscaling data formats for deep learning. *arXiv:2310.10537*, 2023.
- [6] W. Kwon et al. Efficient memory management for large language model serving with PagedAttention. *SOSP*, 2023.
- [7] Z. Ye et al. FlashInfer: Efficient and customizable attention engine for LLM inference serving. *arXiv:2501.01005*, 2025.
- [8] L. Zheng et al. SGLang: Efficient execution of structured language model programs. *arXiv:2312.07104*, 2023.
- [9] A. Kerr et al. CUTLASS: Fast linear algebra in CUDA C++. NVIDIA, 2017–2025. <https://github.com/NVIDIA/cutlass>.